

Annotation Guidelines for Explainability Annotations for Legal Judgment Prediction in Switzerland

Nina Baumgartner

1 Introduction

1.1 Annotation Goal

Recently [Niklaus, Chalkidis, and Stürmer \(2021\)](#) presented a dataset for legal Judgment prediction including 85k Swiss Federal Supreme Court decisions. Using Hierarchical BERT they achieved a Macro-F1 Score of approximately 68-70%. Nevertheless, the inner workings of such models are still mostly unknown and the results thus not interpretable. For this reason, this annotation task focuses on the explainability of these predictions. With your annotation you will give your insight as a legal expert and tag parts of the facts that support or oppose the judgment. These guidelines are based on the work of [Reiter \(2020\)](#), [Leitner, Rehm, and Moreno-Schneider \(2019\)](#) and [Pustejovsky and Stubbs \(2012\)](#). They are a work-in-progress in collaboration with Lynn Grau, Angela Stefanelli and Thomas Lüthi.

1.2 Dataset

For this task a subset of the Swiss court ruling corpus presented by [Niklaus et al. \(2021\)](#) containing 108 cases was created. In most of these cases the model worked well and indicated the correct judgment form the facts given to it. The 108 cases are equally distributed among the three languages German, French and Italian. Each language set contains six cases over six years. With each year having two cases per legal area¹: One with the verdict approved and one with the verdict dismissed.

1.3 Disclaimer

This document is a work-in-progress. If you have questions or find any errors in these instructions while doing the annotation please feel free to contact the maintainer. Please help with collecting examples to complete these guidelines.

2 Annotation Entities

Although you will only be annotating the fact section of a ruling, you will have access to the full document (via a link on Prodigy) and the judgment will be clearly indicated on the prodigy interface. You can and should use these other resources as an indicator on which part of the facts are of greatest importance.

2.1 Sentences and Sub-Sentences

The annotation will focus on the sentences and sub-sentences. For this task we define a sentence as a self-contained linguistic unit consisting of multiple words, terminated with a period, semicolon, colon, question mark or exclamation mark. A entire sentence is the largest entity to be annotated. A sentence can consist of multiple sub-sentences usually separated with a "and" or a comma. It is possible that a sentence contains two sub-sentences opposing each other, which should be consequently annotated with different labels. These sub-sentences are the smallest units that should be annotated. We hope that by choosing this scope it is possible to indicate what the different parts of the sentences denote in the context of the judgment and to subsequently better explain the decisions of the model.

¹The chosen legal areas are categorized as penal law, social law and civil law

2.2 Lower Court

In addition to sentences you will also have to annotate the last lower court of each case. As seen in Figure 1 the Rubrum of the ruling indicates the last lower court. The last lower court is composed of the name of the court e.g. "Verwaltungsgericht" and the location "Kanton Luzern".

Bundesgericht
Tribunal fédéral
Tribunale federale
Tribunal federal

{T 0/2}
8C_892/2009

Urteil vom 15. April 2010
I. sozialrechtliche Abteilung

Besetzung
Bundesrichter Ursprung, Präsident,
Bundesrichterinnen Leuzinger, Niquille,
Gerichtsschreiber Grunder.

Verfahrensbeteiligte
D._____, vertreten durch
Rechtsanwalt Dr. Bruno Häfliger,
Beschwerdeführerin,

gegen

IV-Stelle Luzern, Landenbergstrasse 35,
6005 Luzern, Beschwerdegegnerin.

Gegenstand
Invalidenversicherung (Invalidenrente; Revision),

Beschwerde gegen den Entscheid des Verwaltungsgerichts des Kantons Luzern
vom 16. September 2009.

Figure 1: Example of a Rubrum with the lower court highlighted.

3 Annotation Categories

To annotate the sentences of each fact section you will be using two labels, *Supports judgment* and *Opposes verdict*. You should also highlight the lower court for each judgement. In addition, you will be given several options for dealing with problematic cases, which should help to improve the dataset, these guidelines and the annotations themselves.

3.1 Supports Judgment

This label is used when a sentence or sub-sentence supports the judgment. Every sub-sentence that supports the judgment should be annotated. @ToDo Example and more explanation

3.2 Opposes Judgment

This label is used when a sentence or sub-sentence opposes the judgment. Every sub-sentence that opposes the judgment should be annotated. @ToDo Example and more explanation

3.3 Lower Court

This label is used to highlight the last lower court of the case. To label the last lower court highlight the name and the location of the court as one instance (see Figure 2).

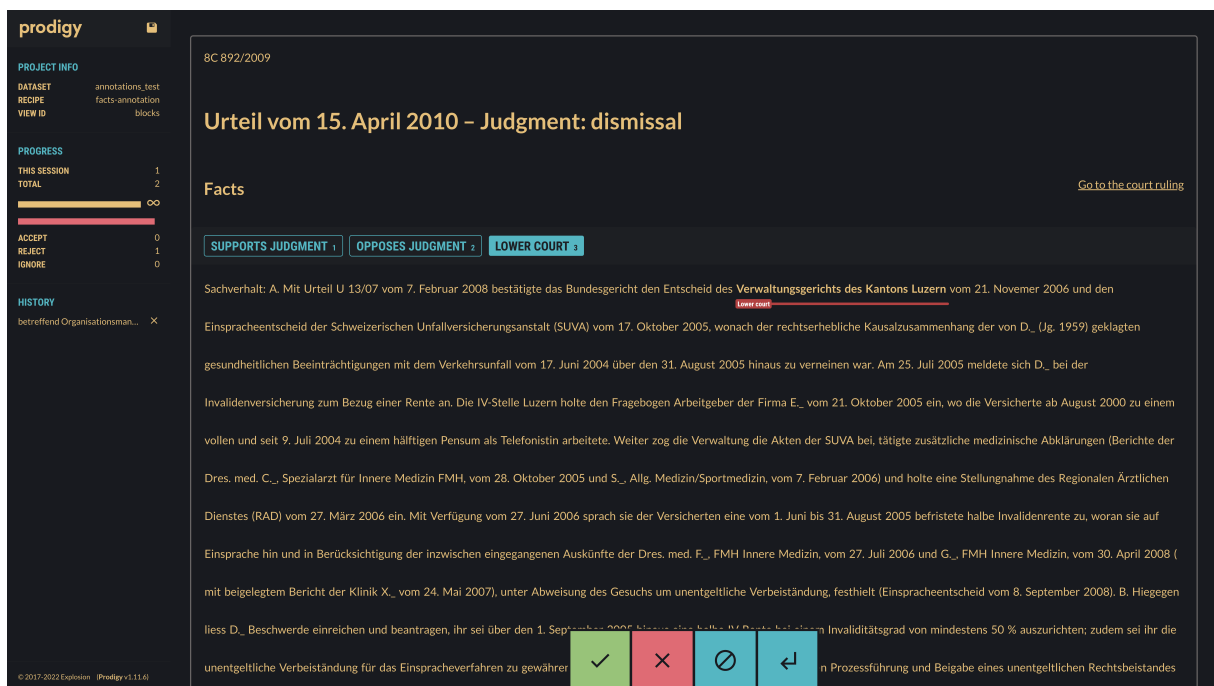


Figure 2: Example of a highlighted lower court in Prodigy.

3.4 Neutral

Every not-labeled sub-sentence is considered neutral. This is not a label per se but merely how the system interprets words or sentences which are not assigned one of the labels above. It is important for the analysis that even the neutral sentences are annotated which in our case means to omit them.

One example in German of a neutral expression which should not be tagged with a label is the word *"Sachverhalt:"*. This word only indicates the beginning of the fact section and should be left out as a neutral part of the facts because it does not give us any further information on the explainability of the judgment.

3.5 Problematic Cases

Problematic cases can occur because this annotation task is an iterative process. For now, we differentiate between three possible types of such cases.

3.5.1 Rejected Cases

If a case is badly tokenized² or there is another formal error it should be rejected. Please state your reasoning in the comment window using the comment pattern below and reference the [Reject or Ignore a Case](#) section of this document for the details on how to properly reject a case. Figure 3 is an example of a case with formal errors.

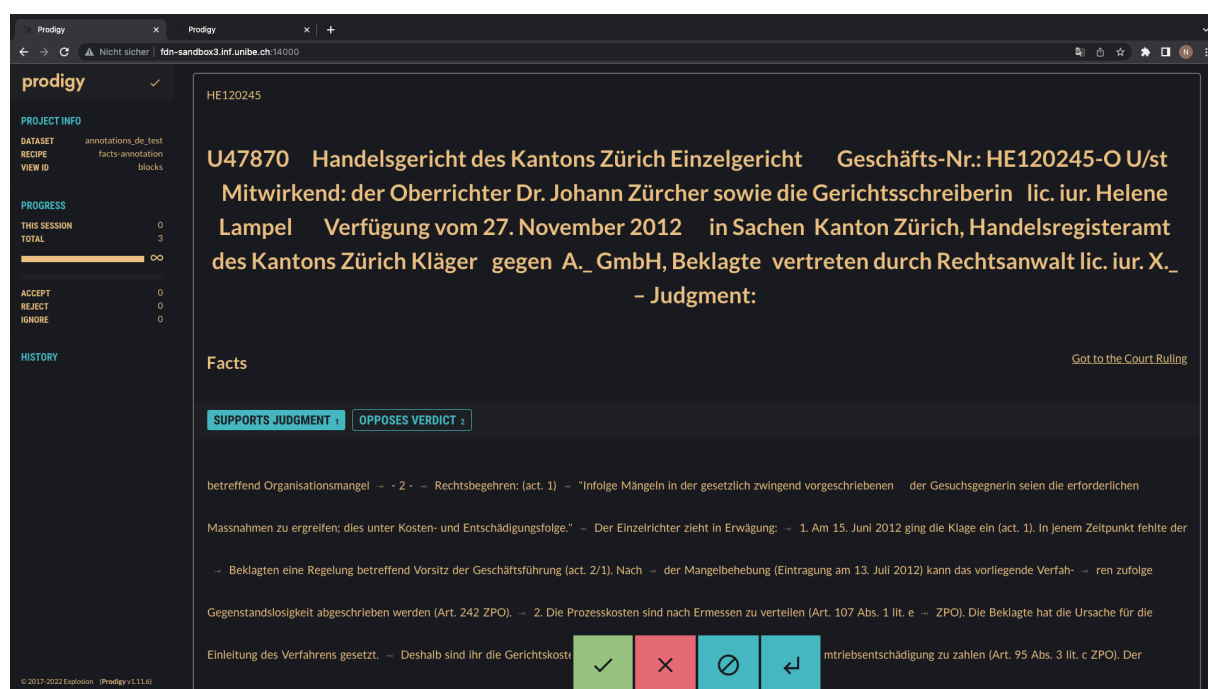


Figure 3: Example of a case containing a formal error which should be rejected. Here the title was parsed incorrectly, the judgment is missing and the facts are tokenized wrongly and incomplete.

3.5.2 Ignored Cases

If a case is too short or otherwise unfit for the annotation it should be ignored. To ignore it please state your reasoning in the comment section and follow the steps explained in the [Reject or Ignore a Case](#) section of this document below.

²Tokenized means that the system did not properly separate the words.

3.5.3 Other Problematic Cases

There might be cases without formal errors where you have difficulties to annotate (neither reject nor ignore). In such cases, please annotate to the best of your ability and explain your reasoning in the comment section.

3.5.4 Comment Structure

Comment for rejecting and ignoring case

Number of case – Annotators name

- Why did you ignore/reject this case?

Comment for generally problematic case

Number of case – Annotators name

- Why is this case problematic and difficult to annotate?
- How did you decide on your annotation?

4 Implementation: How to Annotate the Dataset Using Prodigy

This section explains how to use the annotation tool Prodigy³. We built a custom recipe for this task which lets you annotate the facts section of a given court decision.

4.1 Access

The Prodigy instance can only be accessed via the University of Bern network. If you want to annotate from home you must use the VPN of the University of Bern⁴. If you are connected to the university network you can access Prodigy via one of the following URLs:

- German case annotations:
 - Angela: <http://fdn-sandbox3.inf.unibe.ch:11000/?session=angela>
 - Lynn: <http://fdn-sandbox3.inf.unibe.ch:11000/?session=lynn>
 - Thomas: <http://fdn-sandbox3.inf.unibe.ch:11000/?session=thomas>
- French case annotations: <http://fdn-sandbox3.inf.unibe.ch:12000/>
- Italian case annotations: <http://fdn-sandbox3.inf.unibe.ch:13000/>

Before you can start you will be asked to provide a *username* and a *password*, which will be given to you by the maintainer of the annotation process. After the login procedure you should now see an overview of the case and you can start with your annotation.

4.2 Annotate a Sub-Sentence

To label a phrase with a tag, highlight it with your cursor and choose the corresponding label. To delete a tag simply click on the tagged words again. As seen in Figure 4 the two labels appear in two different colors. By hovering over an annotated section the delete toggle appears.

If you are happy with your annotation you can accept it by clicking on the green check labeled with [1] in Figure 5 and save it by pressing the save button in the left corner referenced by the number [2]. To see your progress you can look at the information displayed on the left (see number [3] on Figure 5). If you want to access the original document you can click on the link in the right corner (see number [4]). Please do not forget to save your progress using the save button [2].

If you want to skip a case, because you already annotated it. Please use the accept button [1] to get to the next case.

³<https://prodi.gy/>

⁴https://serviceportal.unibe.ch/sp?id=kb_article_views&id=00cb11e51b005050134ddc6a9b4bcb49




Figure 4: Screenshot of sentence labeling in prodigy.

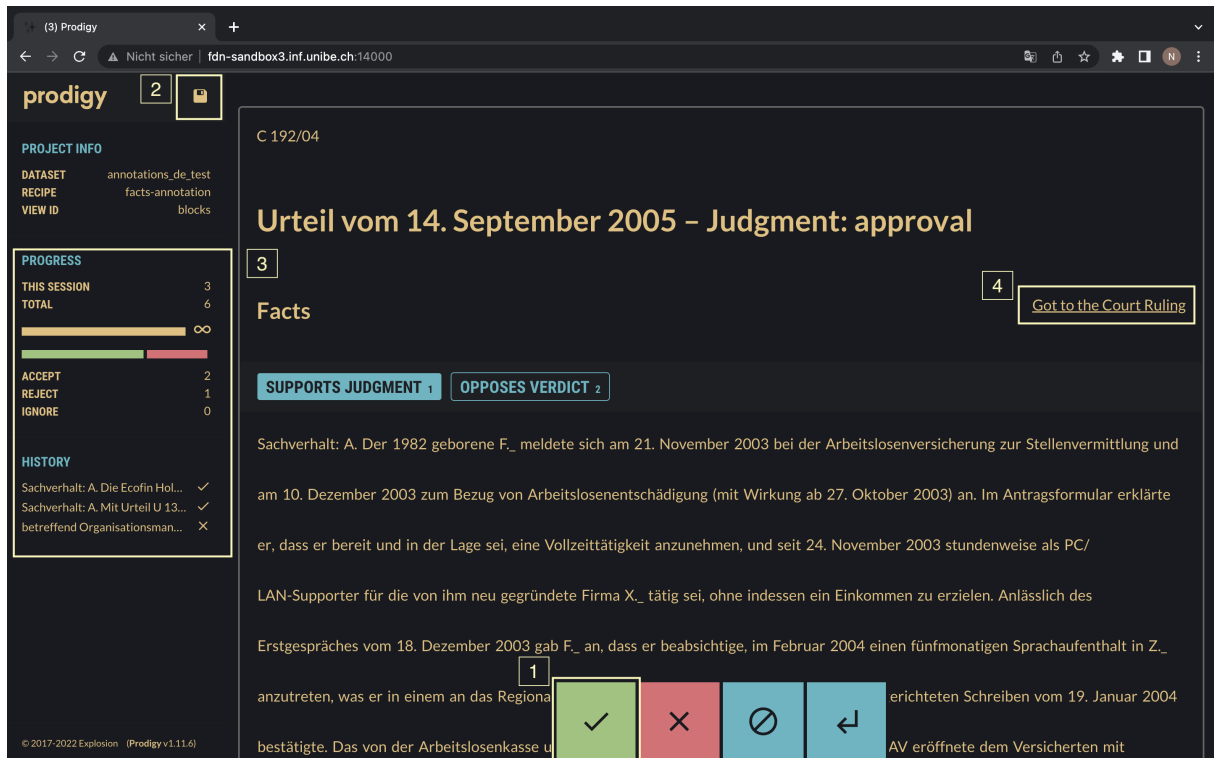


Figure 5: Screenshot of the case overview on prodigy

4.3 Reject or Ignore a Case

To reject a case state your reasoning in the comment section and press the red cross to reject it. To ignore it, press the blue button with the stop signal after commenting. Do not forget to save your progress. Figure 6 shows the interface of the comment section and the ignore and reject buttons.

5 Change Log

This change log documents the progress of these guidelines. When adapting these guideline please also add a new entry to the changelog using the following structure

Template

Date – Title of changes

- Which parts were changed in which iteration?
- Why was this part changed

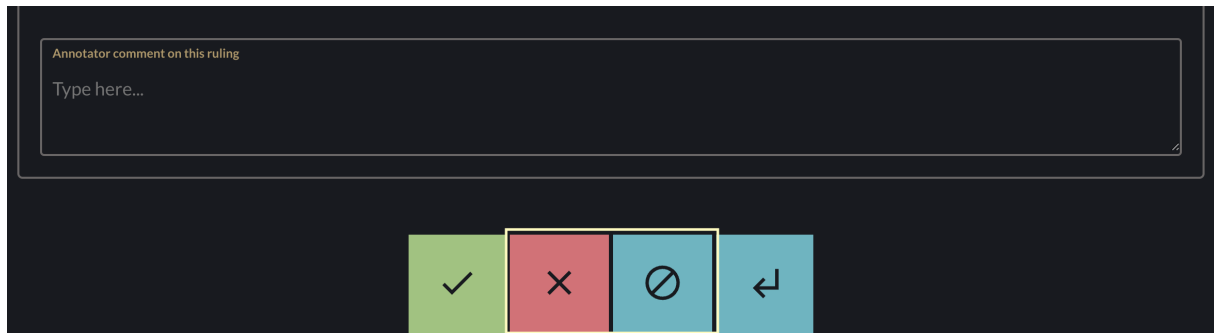


Figure 6: Reject and Ignore buttons

Change log

10.04.2022 – Formal changes after first feedback

- Which parts were changed in which iteration?
 - Changed E. Leitner reference to the published article.
 - Corrected some spelling errors.
 - Integrated the figures into the text of the guidelines.
 - Changed label "opposes verdict" to "opposes judgment"

- Why was this part changed

With this first adaption of the guidelines we mainly worked on some formal errors to standardize the format and clarify the instruction (especially with integrating the figures into the text). The label was changed to make the annotation and their interpretation more consistent.

23.04.2022 – Changes to Prodigy setup and new label

- Which parts were changed in which iteration?
 - Named multi-user sessions were added to the Prodigy setup, which changed the annotators URLs in this document.
 - The label lower court was added as a new annotation category and subsequently to the prodigy setup. Explanation how and when to use it were added to section 2 and 3.
 - Directions on how to skip already annotated cases were added. To section 4.2

- Why was this part changed

The named multi-user session was a pending part of the prodigy setup, which is now resolved. The URLs of the annotators had to be adapted accordingly. After a meeting with the lawyer and annotator Thomas Lüthi, we decided on adding the new label "lower court", to highlight it as a separate entity additionally to the existing two labels. Correct sessions were not yet implemented in the first setup of Prodigy used for some annotations, for this reason directions on how to skip a case a already annotated case were added.

References

- Leitner, E., Rehm, G., & Moreno-Schneider, J. (2019, 9). Fine-grained Named Entity Recognition in Legal Documents. In M. Acosta, P. Cudré-Mauroux, M. Maleshkova, T. Pellegrini, H. Sack, & Y. Sure-Vetter (Eds.), *Semantic systems. the power of ai and knowledge graphs. proceedings of the 15th international conference (semantics 2019)* (pp. 272–287). Karlsruhe, Germany: Springer. (10/11 September 2019)
- Niklaus, J., Chalkidis, I., & Stürmer, M. (2021). *Swiss-judgment-prediction: A multilingual legal judgment prediction benchmark*. arXiv. Retrieved from <https://arxiv.org/abs/2110.00806>

DOI: 10.48550/ARXIV.2110.00806

- Pustejovsky, J., & Stubbs, A. (2012). *Natural language annotation for machine learning* (Nos. Bd. 9,S. 878). O'Reilly Media, Incorporated.
- Reiter, N. (2020). Anleitung zur erstellung von annotationsrichtlinien. In N. Reiter, A. Pichler, & J. Kuhn (Eds.), *Reflektierte algorithmische textanalyse: Interdisziplinäre(s) arbeiten in der creta-werkstatt* (pp. 193–202). De Gruyter. Retrieved from <https://doi.org/10.1515/9783110693973-009>
DOI: doi:10.1515/9783110693973-009